

Anuraj Singh

+91 9838242460
anurajsingh0909@gmail.com
[LinkedIn](#) · [GitHub](#) · [Portfolio](#)

Senior Data Engineer with 5+ years designing and operating large-scale Spark ETL, batch and streaming pipelines, and cloud data platforms (Spark, Scala, SQL, Python, Airflow, AWS). Specialize in high-volume processing, data quality, performance tuning, and AI-ready / analytics datasets for product, advertising, and ML use cases.

Professional Experience

Senior Data Engineer (Full-Time) | JioHotstar | Remote Sep 2026 – Present

- Build and maintain large-scale ETL pipelines for AdTech advertising insights and analytics; develop and optimize distributed Spark applications (Scala / PySpark) and SQL for batch and streaming workloads.
- Drive data quality, performance, and cost efficiency across production pipelines; partner with engineering, product, and analytics on cloud and containerized data platform deliveries.

Senior Software Engineer - AI Evaluation & Benchmarks (Contract) | Handshake AI | Remote Jul 2026 – Sep 2026

- Developed Python-based AI evaluation pipelines for automated benchmark execution, LLM quality assessment, and deterministic validation using reproducible evaluation workflows.
- Built benchmark datasets, automated evaluation and reporting pipelines, and improved engineering workflows using Docker, Git, Linux, and production software engineering practices.

AI Data Engineer (Contract) | Mercor | Remote Aug 2025 – Jul 2026

- Architected scalable Python, SQL, and AWS-based data pipelines to transform and standardize 100K–500K+ records into AI-ready datasets for LLM training, model evaluation, and analytics workloads.
- Delivered AI-ready data platforms supporting machine learning and Generative AI workloads, achieving 99.5%+ data accuracy through robust quality assurance and scalable data engineering practices.

Data Engineer (Full-Time) | Team Computers | Remote Aug 2024 – Jul 2025

- Built and maintained scalable ETL/ELT pipelines using Python, SQL, Apache Spark (PySpark), and Apache Airflow to process high-volume enterprise data, improving pipeline performance and reducing manual data processing by 50%.
- Designed and optimized PostgreSQL/MySQL data models while collaborating with cross-functional teams to automate data workflows, implement data quality validation, and maintain production pipelines with 99.9% reliability for business-critical reporting.

Project Engineer (Full-Time) | Wipro | Remote (Pune) Sep 2021 – Aug 2024

- Optimized enterprise data pipelines and workflows using Python, SQL, and AWS, integrating HPSM and OSM systems to resolve 100+ data-related issues daily while improving data reliability and contributing to a CSAT score of 4.8/5.0.
- Modernized Python- and SQL-based data transformation, reporting, and automation solutions using AWS, Tableau, and Excel, reducing manual reporting effort by 25% and improving business visibility through analytics-driven insights.

Key Skills

Data Engineering & Pipelines: Apache Spark (Scala, PySpark, Spark SQL), ETL/ELT, Apache Airflow, Kafka, AWS Glue, dbt, Delta Lake

Programming Languages: Python, Scala, SQL

Databases & Data Management: PostgreSQL, MySQL, MongoDB, Snowflake, Redis

Cloud & Platforms: AWS (S3, Lambda, Redshift, IAM, EC2), GCS, Kubernetes

Data Processing & Quality: Pandas, NumPy, Great Expectations, Parquet, Avro

APIs & Backend: FastAPI, REST APIs

DevOps & Tooling: Docker, Git, GitHub, GitHub Actions, Linux, Jenkins

AI & GenAI: LLM Data Pipelines, AI-ready datasets, LangChain, OpenAI API, RAG Pipelines